

Intro to OpenMP

Subodh Sharma

Indian Institute of Technology Delhi



An API Standard

- A directive-based method to
 - invoke parallel computations on share-memory multiprocessors
- Specified for C/C++ and Fortran.
- Represents **fork-join** model of parallelism
- Jointly developed by all major h/w and s/w vendors:
 - HP, Fujitsu, AMD, IBM, Intel, Nvidia, Cray, TI, Oracle ..

Motivations?

- To primary parallelize regular loops.
 - Such as those in matrix multiplication
- Now supports task-parallelism too
 - Inspired from Cilk, TBBs, X10, Chapel
- Latest version has support for:
 - Accelerators
 - Atomics
 - Thread affinity
 - User defined reduction

How to Compile OpenMP

- `>= gcc4.2` supports OpenMP 3.0
 - `gcc -fopenmp example.c`
- To change the number of threads:
 - `setenv OMP_NUM_THREADS 4 (tcsh)` or `export OMP_NUM_THREADS=4(bash)`
 - can also be provided in the code it self.

```
/*Introducing omp parallel */
```

```
#include <omp.h> // required header to write OpenMP code  
#include <stdio.h>
```

```
int main (){
```

```
}
```

```
/*Introducing omp parallel */
```

```
#include <omp.h> // required header to write OpenMP code  
#include <stdio.h>
```

```
int main (){
```

```
// sentinel directive_name [clause[,clause]...]
```

```
#pragma omp parallel /* compiler directive */
```

```
{  
    int tid = omp_get_thread_num(); /* Library call */  
    printf("hello world %d \n", tid);
```

```
} // implicit barrier synchronization here!
```

```
}
```

Hello world 0

Hello world 2

Hello world 1

Hello world 3


Why 4 threads? Ans: Default – can be changed by
suitably setting the environment var: **OMP_NUM_THREADS**

```
    /*Introducing omp parallel num_threads(x) */  
  
#include <omp.h> // required header to write OpenMP code  
#include <stdio.h>  
  
int main ()  
  
#pragma omp parallel num_threads (2)  
{  
    int tid = omp_get_thread_num();  
    printf("hello world %d \n", tid);  
  
} // implicit barrier synchronization here!  
  
}
```

Hello world 0
Hello world 1

Execution Model

[courtesy: xyuan, FSU]



- Execution model host-centric – host device offloads targets to target devices
- Threads can't migrate from one device to other
- A team is created when a parallel region is encountered

Execution Model

- **parallel** region creates tasks and map to threads
- implicit barrier at the end of **parallel** region
- Only the master resumes beyond **parallel** region
- Threads in a **team** executes **worksharing** tasks cooperatively

Memory Model

- OpenMP API provides a **relaxed-consistency**, shared-memory model
- Threads is allowed to have its own *temporary view* of memory.
 - cache, local storage (thread-private mem), registers
- Data-sharing attribute: **shared, private**
 - For each **private** var new

Memory Model

- a thread's temporary view of memory is not required to be consistent with memory at all times
 - **Solution:** **Flush** operation
- **Flush:** Strong, Rel, Acq
 - Strong: thread writes multiple times between two strong flush ops, guarantees the last write to be in memory.

```
/*Introducing omp lib functions */
```

```
#include <omp.h> // required header to write OpenMP code  
#include <stdio.h>
```

```
int main (){
```

```
#pragma omp parallel num_threads (4)  
{
```

```
    int numt = omp_get_num_threads();  
    int tid = omp_get_thread_num();  
    printf("hello world %d of %d \n", tid, numt);
```

```
} // implicit barrier synchronization here!
```

```
}
```

Hello world 0 of 4

Hello world 2 of 4

Hello world 3 of 4

Hello world 1 of 4

*/*Exposing data race*/*

```
int main (){  
  
    int numt, tid ;  
  
    #pragma omp parallel num_threads (4)  
    {  
        numt = omp_get_num_threads();  
        tid = omp_get_thread_num();  
        printf("hello world %d of %d \n", tid, numt);  
  
    } // implicit barrier synchronization here!  
  
}
```

Hello world 0 of 4
Hello world 2 of 4
Hello world 3 of 4
Hello world 1 of 4

Data race not exposed!

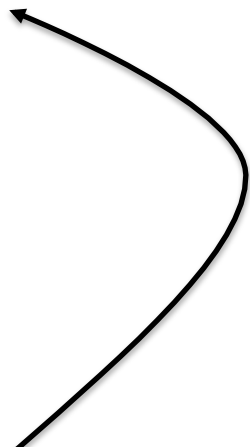
*/*Exposing data race*/*

```
int main (){  
  
int numt, tid ;  
  
#pragma omp parallel num_threads (4)  
{  
    numt = omp_get_num_threads();  
    tid = omp_get_thread_num();sleep(1);  
    printf("hello world %d of %d \n", tid, numt);  
  
} // implicit barrier synchronization here!  
  
}
```

Hello world 3 of 4
Hello world 3 of 4
Hello world 3 of 4
Hello world 3 of 4

*/*Fixing data race*/*

```
int main (){  
  
int numt, tid ;  
  
#pragma omp parallel num_threads (4) shared (numt) private(tid)  
{  
    numt = omp_get_num_threads();  
    tid = omp_get_thread_num();sleep(1);  
    printf("hello world %d of %d \n", tid, numt);  
} // implicit barrier synchronization here!  
}
```



Hello world 1 of 4
Hello world 0 of 4
Hello world 2 of 4
Hello world 3 of 4

tid becomes a thread local
variable – put on thread stack!

*/*Thread local var storage*/*

```
int main (){
```

```
int numt, tid ;
```

```
printf("hello world %d of %d: A(numt) = %x, A(tid) = %x  \n", tid, numt);
```

```
#pragma omp parallel shared (numt) private(tid)
```

```
{
```

```
    numt = omp_get_num_threads();
```

```
    tid = omp_get_thread_num();
```

```
    printf("hello world %d of %d: A(numt) = %x, A(tid) = %x  \n", tid, numt,  
&numt, &tid);
```

```
}
```

```
}
```

tid for Master thread gets duplicated

from MASTER: A(numt) = 726fd678, A(tid) = 726fd67c

Hello world 0 of 4: A(numt) = 726fd678, A(tid) = 726fd64c

Hello world 3 of 4: A(numt) = 726fd678, A(tid) = fa87ae5c

Hello world 1 of 4: A(numt) = 726fd678, A(tid) = fb91ce5c

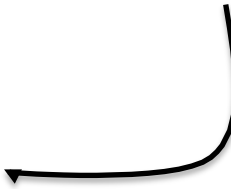
Hello world 2 of 4: A(numt) = 726fd678, A(tid) = fb0cbe5c



*/*only one thread populating numt*/*

```
int main (){  
    int numt, tid ;  
  
    numt = omp_get_num_threads();  
  
    #pragma omp parallel shared (numt) private(tid)  
    {  
        tid = omp_get_thread_num();  
        printf("hello world %d of %d\n", tid, numt);  
    }  
}
```

Wrong move – threads are not even created yet!!



Hello world 3 of 1
Hello world 2 of 1
Hello world 0 of 1
Hello world 1 of 1

*/*only one thread populating numt*/*

```
int main (){
```

```
int numt, tid ;
```

```
#pragma omp parallel private(tid) {
```

```
tid = omp_get_thread_num();
```

```
if (tid == 0)
```

```
    numt = omp_get_num_threads();
```

```
}
```

```
#pragma omp parallel shared (numt) private(tid)
```

```
{
```

```
    tid = omp_get_thread_num();
```

```
    printf("hello world %d of %d\n", tid, numt);
```

```
}
```

```
}
```

Hello world 0 of 4

Hello world 1 of 4

Hello world 3 of 4

Hello world 2 of 4

computing tid multiple times



/*only one thread populating numt & use of barrier*/

```
int main (){  
    int numt, tid ;  
  
    #pragma omp parallel shared (numt) private(tid)  
    {  
        tid = omp_get_thread_num();  
  
        if(tid == 0) numt = omp_get_num_threads();  
  
        #pragma omp barrier  
  
        printf("hello world %d of %d\n", tid, numt);  
    }  
}
```

Hello world 0 of 4
Hello world 1 of 4
Hello world 3 of 4
Hello world 2 of 4

*/*threadprivate(vars) – buggy usage*/*

```
int tvar;
```

```
#pragma omp threadprivate (tvar)
```

```
int main (){
```

```
int numt, tvar = 10;
```

```
printf("From MASTER: A(tid) = %x \n", &tid);
```

```
#pragma omp parallel shared (numt) private(tid)
```

```
{
```

```
    tid = omp_get_thread_num();
```

```
    numt = omp_get_num_threads();
```

```
    printf("hello world %d of %d: tvar = %d, A(tvar) = %x \n", tid, numt, +  
+tvar, &tvar);
```

```
}
```

```
}
```

Hello world 0 of 4: tvar = 11, A(tvar) = 880d0828

Hello world 2 of 4: tvar = 12, A(tvar) = 880d0828

Hello world 3 of 4: tvar = 14, A(tvar) = 880d0828

Hello world 1 of 4: tvar = 13, A(tvar) = 880d0828

`/*threadprivate(vars) – copyin usage*/`

```
int tvar;
```

```
#pragma omp threadprivate (tvar)
```

```
int main (){
```

```
int numt;
```

```
tvar = 10;
```

```
printf("From MASTER: A(tid) = %x \n", &tid);
```

```
#pragma omp parallel shared (numt) private(tid) copyin(tvar)
{
```

```
    tid = omp_get_thread_num();
```

```
    numt = omp_get_num_threads();
```

```
    printf("hello world %d of %d: tvar = %d, A(tvar) = %x \n", tid, numt, +  
+tvar, &tvar);
```

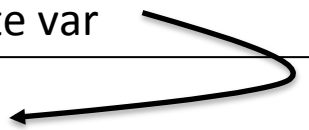
```
}    Hello world 0 of 4: tvar = 11, A(tvar) = 880d0828
```

```
}    Hello world 2 of 4: tvar = 11, A(tvar) = 6f1d27b8
```

```
    Hello world 3 of 4: tvar = 11, A(tvar) = 6d73e6f8
```

```
    Hello world 1 of 4: tvar = 11, A(tvar) = 6df8f6f8
```

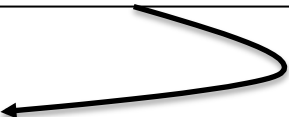
copies the value of master's
threadprivate var



`/*private(vars) – firstprivateusage*/`

```
int main (){  
  
    int numt;  
    int tvar = 10;  
  
    printf("From MASTER: A(tid) = %x \n", &tid);  
  
    #pragma omp parallel shared (numt) private(tid) firstprivate(tvar)  
    {  
        tid = omp_get_thread_num();  
  
        numt = omp_get_num_threads();  
  
        printf("hello world %d of %d: tvar = %d, A(tvar) = %x \n", tid, numt, +  
+tvar, &tvar);  
  
    }  
}
```

initialize the value with prior values
available before the region



Hello world 0 of 4: tvar = 11, A(tvar) = 880d0828
Hello world 2 of 4: tvar = 11, A(tvar) = 6f1d27b8
Hello world 3 of 4: tvar = 11, A(tvar) = 6d73e6f8
Hello world 1 of 4: tvar = 11, A(tvar) = 6df8f6f8

*/*only one thread populating numt & use of thread single*/*

```
int main (){  
    int numt, tid ;  
  
    #pragma omp parallel shared (numt) private(tid)  
    {  
        tid = omp_get_thread_num();  
  
        #pragma omp single  
        {  
            numt = omp_get_num_threads();  
        } // implicit barrier  
  
        printf("hello world %d of %d\n", tid, numt);  
    }  
}
```

Hello world 0 of 4
Hello world 1 of 4
Hello world 3 of 4
Hello world 2 of 4

*/*use of copyprivate with single directive*/*

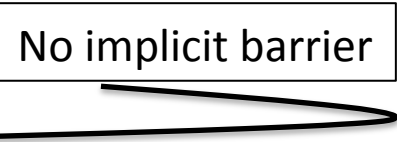
```
int main (){  
    int numt, tid, a ;  
  
    #pragma omp parallel shared (numt) private(tid, a)  
    {  
        tid = omp_get_thread_num();  
  
        #pragma omp single copyprivate(a)  
        {  
            numt = omp_get_num_threads();  
            a = tid;  
        } // implicit barrier  
  
        printf("hello world %d of %d\n: a = %d", tid, numt, a);  
    }  
}
```

broadcasts a thread private var value to others

Hello world 0 of 4:a = 0
Hello world 1 of 4:a = 0
Hello world 3 of 4:a = 0
Hello world 2 of 4:a = 0

/*use of thread single with nowait*/

```
int main (){  
    int numt, tid ;  
  
    #pragma omp parallel shared (numt) private(tid)  
    {  
        tid = omp_get_thread_num();  
        #pragma omp single nowait  
        {  
            numt = omp_get_num_threads();  
        }  
  
        printf("hello world %d of %d\n", tid, numt);  
    }  
}
```



Hello world 0 of 4
Hello world 1 of 4
Hello world 3 of 4
Hello world 2 of 0

*/*use of thread single with only master thread*/*

```
int main (){  
    int numt, tid ;  
  
    #pragma omp parallel shared (numt) private(tid)  
    {  
        tid = omp_get_thread_num();  
  
        #pragma omp master  
        {  
            numt = omp_get_num_threads();  
        } // no implicit barrier!  
  
        printf("hello world %d of %d\n", tid, numt);  
    }  
}
```

Only master thread executes

Hello world 0 of 4
Hello world 1 of 4
Hello world 3 of 4
Hello world 2 of 0

Computing Pi via Integral of $4.0/(1.0 + x*x)$

/*Synchronization in OpenMP*/

/*Barriers*/

```
int main (){
```

```
int numt, tid ;
```

```
#pragma omp parallel shared (numt) private(tid)
{
```

```
    tid = omp_get_thread_num();
```

```
    if(tid == 0) numt = omp_get_num_threads();
```

```
#pragma omp barrier
```

```
    printf("hello world %d of %d\n", tid, numt);
```

```
}
```

```
}
```

`/*Synchronization in OpenMP*/`

`/*Locks*/`

`omp_init_lock(omp_lock_t *)`

- Nestable locks are declared with the type

`omp_nest_lock_t`

`omp_set_lock()` -- acquires lock

`omp_unset_lock()` - releases lock

`omp_destroy_lock()` - free memory

`omp_test_lock ()` - set the lock if available, else
return w/o blocking

```
while(!flag) {  
    flag = omp_test_lock();
```

```
}
```

/*Synchronization in OpenMP*/

/*Ordered Construct*/

```
int main() {
```

```
int i;
```

```
# pragma omp parallel for ordered
```

```
{
```

```
    for (i=0; i < 5; i++)
```

```
        foo();
```

```
}
```

```
}
```

/*Synchronization in OpenMP*/

/*Critical Construct*/

sum = 0;

pragma omp parallel shared (n,a,sum) private(tid,
sumLocal)

{

tid = omp_get_thread_num();
sumLocal = 0;

#pragma omp for
for(i=0;i<n;i++)
sumLocal += a[i];

#pragma omp critical (update_sum) // name of the block
sum += sumLocal;

}
}

/*Synchronization in OpenMP*/

/*atomic Construct*/

sum = 0;

pragma omp parallel shared (n,a,sum) private(tid,
sumLocal)

{

tid = omp_get_thread_num();

sumLocal = 0;

#pragma omp for

for(i=0;i<n;i++)

sumLocal += a[i];

#pragma omp atomic

sum += sumLocal + foo();

}

}

*/*use of Loop work sharing*/*

```
#pragma omp parallel
{
#pragma omp for
    for (i = 0; i < N; i++)
        do_stuff(); // a[i] += b[i]
} // implicit barrier synchronization here!
```

- Limited to loops where iterations are a-priori known!
- The loop iteration variable is treated as thread local by the compiler
- iteration-to-threads mapping compiler defined!
- Loops must be free from loop-carried dependencies!
- clauses supported: **private, firstprivate, reduction, schedule, nowait, ..**

*/*use of Loop work sharing*/*

```
#pragma omp parallel for  
{  
    for (i = 0; i < N; i++)  
        do_stuff(); // a[i] += b[i]  
}
```

Succinct form when
there is only a loop
to be parallelized in
“omp parallel” region

*/*cyclic distribution of the same workload*/*

```
#pragma omp parallel  
{  
    int id, Nthrds, istart, iend;  
    id = omp_get_thread_num();  
    Nthrds = omp_get_num_threads();  
    // figure out istart, iend  
    for (i = istart; i < iend; i++)  
        do_stuff(); // a[i] += b[i]  
}
```

/*use of Loop work sharing with schedule clause*/

```
#pragma omp parallel for schedule (static, chunk-size)
{
    for (i = 0; i < N; i++)
        do_stuff(); // a[i] += b[i]
}
```

- granularity of workload-per-thread is **chunk-size**
 - **chunk-size** – is a continuous non-empty subset of iteration space.
- Least overhead; chunks are assigned to threads in round-robin manner in the order of the thread IDs.
 - The last thread may get smaller number of iterations
- each thread gets at most 1 chunk if the chunk-size is unspecified

*/*use of Loop work sharing with schedule clause*/*

```
#pragma omp parallel for schedule (dynamic, chunk-size)
{
    for (i = 0; i < N; i++)
        do_stuff(); // a[i] += b[i]
}
```

- iterations are assigned to threads in the team of chunks
 - Each chunk = chunk-size many iterations
- once threads finish with their allotted chunk they request for more
- When no chunk-size specified – the size defaults to 1.
- Mostly used when the workload is unpredictable and poorly balanced

*/*use of Loop work sharing with schedule clause*/*

```
#pragma omp parallel for schedule (guided, chunk-size)
{
    for (i = 0; i < N; i++)
        do_stuff(); // a[i] += b[i]
}
```

- Similar to dynamic scheduling, except that chunks start decreasing over time
 - For chunk-size = 1 --- each chunk size will be proportional to (number of unassigned iterations)/(number of threads in the team)
 - For chunk-size = k --- size of the chunk determined as before but with the restriction that no chunk will have less than k iterations (except possibly for the chunk that contains the last iteration)
 - When no chunk-size is specified, it defaults to 1.

*/*use of Loop work sharing with schedule clause*/*

```
#pragma omp parallel for schedule (runtime)
{
    for (i = 0; i < N; i++)
        do_stuff(); // a[i] += b[i]
}
```

- the decision regarding scheduling is deferred until run time, and the schedule and chunk size are taken from the run-sched-var ICV (Internal Control Variable).

/*Reduction Clause*/

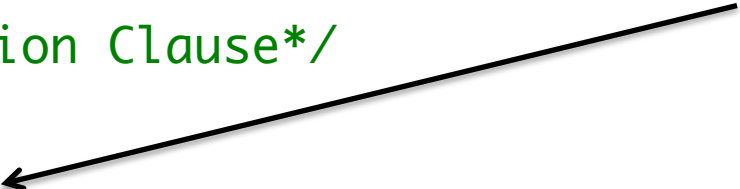
```
#pragma omp parallel for
{
    for (i = 0; i < N; i++)
        sum += a[i];
}
```

Can we do it this way?

Always a good debugging practice

/*Reduction Clause*/

```
#pragma omp parallel for default(none) reduction(+:sum)
{
    for (i = 0; i < N; i++)
        sum += a[i];
}
```



/*Reduction Clause*/

```
#pragma omp parallel for default(none) shared(a,N)
reduction(+:sum)
{
    for (i = 0; i < N; i++)
        sum += a[i];
}
```

/*Reduction Clause explicit implementation*/

```
#pragma omp parallel default(none) shared(a,N,sum)
{
    int sumLocal = 0;
    #pragma omp for
    for (i = 0; i < N; i++)
        sumLocal += a[i];
    #pragma omp critical
    sum += sumLocal
}
```


`/*Reduction Clause*/`

Operators and initial values supported by reduction clause

Operator	Initialization value
+	0
*	1
-	0
&	1
	0
&&	1
	0

*/*Parallel for Collapse*/*

```
void sub(float *a)
{
    int i, j, k;
    #pragma omp for collapse(2) private(i, k, j)
    for (k=kl; k<=ku; k+=ks)
        for (j=jl; j<=ju; j+=js)
            for (i=il; i<=iu; i+=is)
                bar(a,i,j,k);
}
```

The iterations for the k and the j loops are collapsed into one loop – the iteration space is then divided according to the schedule clause.

/*use of Section work sharing*/

```
#pragma omp parallel{  
    #pragma omp sections {  
        #pragma omp section  
        func_x();  
        #pragma omp section  
        func_y();  
    } // sections region has an implicit barrier  
}
```

- each thread executes the region within a section
- each section is executed only once
- if only one thread is there? What is the order of execution of regions?
- if more than two threads?
- Assignment of code blocks to threads is implementation-dependent

/*OpenMP Tasking Constructs*/

```
1  #pragma omp parallel
2  {
3      #pragma omp task
4      Compute1();
5      #pragma omp task
6      Compute2();
7  }
```

Order of execution of
tasks is undefined!



- Binding: a task binds to the innermost enclosing parallel region
- Task synchronization: either by `#pragma omp barrier` or by `#pragma omp taskwait`
 - taskwait: explicit wait on the completion of child tasks

/*OpenMP Task Dependencies*/

```
1  #pragma omp parallel
2  {
3      #pragma omp task depend (OUT:x)
4      x = Compute1();
5      #pragma omp task depend (IN:x)
6      Compute2(x);
7  }
```

- **In Dependency:** generated task will depend on all previously generated sibling tasks with **out** or **inout** clauses
- **Out & Inout Dependency:** generated task will depend on all previously generated sibling tasks with **in**, **out** and **inout** clauses

*/*OpenMP Task Scheduling*/*

```
void Foo(omp_lock_t *lock, int n){  
    for (int i=0 ; i<n ; i++) {  
        #pragma omp task  
        {  
            something_useful();  
            while (!omp_test_lock(lock)){  
                # pragma omp taskyield  
            }  
            something_critical();  
            omp_unset_lock(lock);  
        }  
    }  
}
```

- **taskyield:** the current task can be suspended in favor of the execution of a different task.
 - to avoid deadlocks.

/*Advanced OpenMP features - flush*/

#pragma omp flush

// 1. OpenMP standard specifies that all shared mem modifications are available to all threads at **synchronization points**

//2. Flush forces an update in between sync points

//3. It **does not** synchronize the actions of different threads

//4. Compiler can re-order flush operation – thus, it may not execute exactly at the position relative to other operations as specified by the programmer.

/*Performance Optimization Tips*/

1. Understand memory access patterns and perform **loop interchange**, if necessary

```
for (j=0; j<n; j++) {// Better soln: interchange the loops  
    for (i=0; i<m; i++) {  
        sum += a[i][j];  
    }  
}
```

2. **Loop unrolling** is a powerful technique to avoid loop-overheads

```
for (i=1; i<n; i++) {  
    a[i]= b[i] + 1;  
    c[i] = a[i] + a[i-1] + b[i-1];  
    /* transformed in to:  
    for (i=1; i<n; i+=2) {  
        a[i]= b[i] + 1;  
        c[i] = a[i] + a[i-1] + b[i-1];  
        a[i+1]= b[i+1] + 1;  
        c[i+1] = a[i+1] + a[i] + b[i];  
    }  
}
```


/*Performance Optimization Tips*/

3. Loop with complex access patterns

```
for (j=0; j<n; j++) { // loop tiling: exercise!
    for (i=0; i<m; i++) {

        a[i][j+1]= a[i+1][j] + 1;

    }
}
```

$a[2][2]$ is computed in iter $j=1, i=2$ and used to assign a new value to $a[1][3]$ in iteration $j=2, i=1$

4. Loop fusion/fission: merge/split two loops to create a bigger/smaller loop to increase cache usage.

Eg:

```
for (i=0; i<n; i++) { // loop fission
    c[i]= exp(i/n);
    for (j=0; j<m; j++) {
        a[j][i] = b[j][i] + d[j]*e[i];
    }
}
```

/*Performance Optimization Tips*/

5. Optimize use of barriers;

```
#pragma omp parallel ... {  
  
    #pragma omp for {  
        for (i=0; i<n; i++)  
            a[i] += b[i];  
    }  
  
    #pragma omp for {  
        for (i=0; i<n; i++)  
            c[i] += d[i];  
    }  
  
    #pragma omp for reduction (+:sum)  
        for (i=0; i<n; i++)  
            sum += a[i] + c[i];  
}
```

6. Avoid large critical sections, use atomics where you can

`/*Synchronization*/`

`barrier, critical, atomic, lock`

`/*Work Sharing*/`

`for, sections, single, master`

`/*Clauses to control Work Sharing*/`

`shared, private, firstprivate, copyin, default, nowait,
schedule, reduction`

`/*Additional Library routines, environment variables*/`