

2016CS10367

[illegible]

\ /	\ /													
-----	-----	--	--	--	--	--	--	--	--	--	--	--	--	--

```
#pragma omp parallel for schedule(static)
for(int j=0;j<n;j++){
    for(int i=1;i<n;i++){
        a[i*n+j]+=a[(i-1)*n+j];
    }
}
```

(c)

In foo the critical section and in bar the atomic section i.e line 4 and line 13 ,respectively can lead to data races.

The solution is to

Change the foo code:

```
//locatedinsideaparallelregioninfunctionfoo
#pragma omp atomic
{
    a++;//Modifyaexclusively
}
//locatedinsideaparallelregioninfunctionbar
#pragma omp atomic
{
    b++;//Modifybexclusively
}
#pragma omp atomic
{
    a=a+b;//Modifyaexclusively
}
```

Problem2

(a)

Dependencies:

No data dependency in any loop.

The outer loop is ideally suited for parallelisation than inner loop because if we parallelise the inner loop it will create more overhead due to the creation and joining of threads for all i.

(b)

Code:

Compile:g++ P2.cpp -fopenmp

```
#include <iostream>
#include <vector>
#include <omp.h>
#include <random>
#include <algorithm>
#include <iterator>
#include <functional>
```

```
using namespace std;
```

```
void seqsort(std::vector<unsigned int> &X){
    unsigned int i,j,count,N=X.size();
    std::vector<unsigned int> tmp(N);
    for(i=0;i<N;i++){
        count=0;
        for(j=0;j<N;j++)
            if(X[j]<X[i]||X[j]==X[i]&&j<i)
                count++;
        tmp[count]=X[i];
    }
    std::copy(tmp.begin(),tmp.end(),X.begin());
}
```

```
void parsort(std::vector<unsigned int> &X){
    unsigned int i,j,count,N=X.size();
    std::vector<unsigned int> tmp(N);

    #pragma omp parallel for schedule(static) shared(tmp) private(j,count) num_threads(4)
    for(i=0;i<N;i++){
        count=0;
        for(j=0;j<N;j++)
            if(X[j]<X[i]||X[j]==X[i]&&j<i)
                count++;
        tmp[count]=X[i];
    }
    std::copy(tmp.begin(),tmp.end(),X.begin());
}
```

```
int main(int argc,char const*argv[]){

    random_device rnd_device;
    // Specify the engine and distribution.
    mt19937 mersenne_engine {rnd_device()}; // Generates random integers
```

```

uniform_int_distribution<int> dist {1, 1000};

auto gen = [&dist, &mersenne_engine]() {
    return dist(mersenne_engine);
};

std::vector< unsigned int> vec(10000);
generate(begin(vec), end(vec), gen);
std::vector<unsigned int> vec2=vec;
for (auto i = vec.begin(); i != vec.end(); ++i)
    std::cout << *i << ' ';
std::cout<<"\n";
double start_time = omp_get_wtime();
seqsort(vec);
double time1 = omp_get_wtime() - start_time;
start_time = omp_get_wtime();
parsort(vec2);
double time2 = omp_get_wtime() - start_time;

for (auto i = vec2.begin(); i != vec2.end(); ++i)
    std::cout << *i << ' ';
std::cout<<"\n";
std::cout<<"Time taken by seqsort "<<time1<<endl;
std::cout<<"Time taken by parsort "<<time2<<endl;
return(0);
}

```

Threads	2	4
Speed Up	1.62	1.95
Efficiency	0.81	0.48

Problem3

(a)

The code is actually trying to calculate the sum and uses $\log(n)$ num of steps as the result half of the threads become ideal at each 's' doubling
The problem with the code is the if condition which made most of the thread ideling, instead try to replace it, with
for (unsigned int s=blockDim.x/2; s>0; s=s/2)
{

```

        sdata[tid] += sdata[tid + s];
    __syncthreads();
}

```

According to me __syncthreads cannot be removed as all the thread need to be synchronized till this point.

(b)

Ways to avoid Data Race are:

__syncthreads: act as a barrier

atomicAdd(), atomicSub(), atomicMin(), atomicMax()

Lock, mutex

__global__: this means the function will be called from the host and will run on the device, i.e., the GPU

__device__: Function call from the device and run on the device

__host__: function called from the host and executed on the host

Problem4

Future and Promises provide a way to do asynchronous programming by the means of using signals.

Async creates another thread (depends on parameter may work with a single thread) to launch the function provided in the parameter.

The class template `std::promise` provides a facility to store a value or an exception that is later acquired asynchronously via an `std::future` object created by the `std::promise` object.

Basically, we are creating a communication channel between threads to share the values required by a thread when generated.

Example Code:

Compile:g++ P4.cpp -pthread

```
#include <future>
#include <iostream>
#include <unistd.h>

using namespace std;

int Sum(future<int> &f){
    int sum=0;
    int n=f.get();
    for(int i=1;i<n+1;i++){
        sum+=i;
    }
    cout<<"Result of Thread "<<sum<<endl;
    cout<<"I have thread id "<<this_thread::get_id()<<endl;
    return(sum);
}

int main(int argc, char const *argv[])
{
    int x;
    promise<int> p;//Promises are container of future
    future<int> f=p.get_future();
    future<int> fu=async(launch::async,Sum,ref(f));//Sum function launch asynchronously in a
separate thread(launch::async )
    cout<<"Hello main thread here\n";
    cout<<"I have thread id "<<this_thread::get_id()<<"\n\n\n\n";
    //Do some Work to give the value to Sum function
    sleep(2);//doing some work 2s required (lot of work)
    int val=100;//Ok Now I need the SUM of 1st 100 element
    p.set_value(val);//As promised a value is given to p
    //fu.get() will wait till the time the "Sum " execution is complete,i.e,it is blocking
    x=fu.get();//getting the returned value by thread2

    return 0;
}
```